

(2,4) tree

- ① All the leaf are at same level.
- ② Each node should have atmost three and atleast one key
- ③ " " " atmost four atleast two children

Slide Credit : Yael Moses, IDC Herzliya
Animated Video [B-Tree example](#)



Motivation

- Large differences between time access to disk, cash memory and core memory
- Minimize expensive access (e.g., disk access)
- B-tree: Dynamic sets that is optimized for disks

B-Trees

A B-tree is an M-way search tree with two properties :

1. It is perfectly balanced: every leaf node is at the same depth
2. Every internal node other than the root, is at least half-full, i.e. $M/2-1 \leq \#keys \leq M-1$
3. Every internal node with k keys has $k+1$ non-null children

$$\frac{M}{2}-1 \leq \# \text{key} \leq M-1$$

For simplicity we consider M even and we use $t=M/2$:

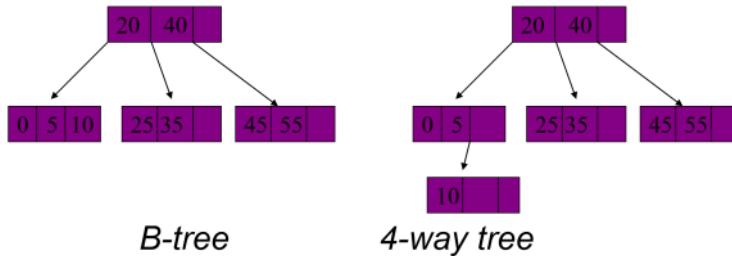
- 2.* Every internal node other than the root is at least half-full, i.e. $t-1 \leq \#keys \leq 2t-1$, $t \leq \#children \leq 2t$

$$\begin{array}{c} [K_1, K_2, \dots, K_r] \\ t-1 \leq r \leq 2t-1 \quad \text{#of keys} \\ t \leq \# \text{children} \leq 2t \end{array}$$

Example: a 4-way B-tree

$$2t-1 = 3$$

$$t = 2$$



B-tree

1. It is perfectly balanced: every leaf node is at the same depth.
2. Every node, except maybe the root, is at least half-full
 $t-1 \leq \#keys \leq 2t-1$
3. Every internal node with k keys has $k+1$ non-null children

B-tree Height

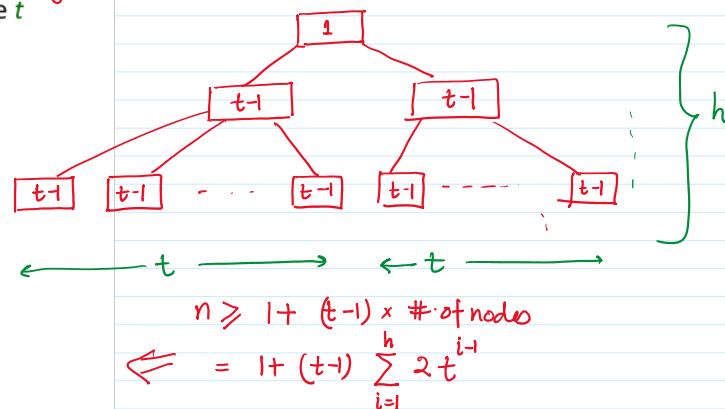
$$t = \# \text{of children for any internal node}$$

Claim: any B-tree with n keys, height h and minimum degree t satisfies:

$$h \leq \log_t \frac{n+1}{2}$$

Proof:

- The minimum number of KEYS for a tree with height h is obtained when:
 - The root contains one key
 - All other nodes contain $t-1$ keys



$$\begin{aligned} n &\geq 1 + (t-1) \times \# \text{of nodes in Btree} \\ &= 1 + (t-1) \times \sum_{i=1}^h 2^{t-1} \end{aligned}$$

$$= 1 + (t-1) \times \sum_{i=1}^h 2^i t^{i-1}$$

$$n \geq 1 + (t-1) \sum_{i=1}^h 2^i t^{i-1}$$

B-Tree: Insert X

- As in M -way tree find the leaf node to which X should be added
- Add X to this node in the appropriate place among the values already there
(there are no subtrees to worry about)
- Number of values in the node after adding the key:
 - Fewer than $2t-1$: done
 - Equal to $2t$: overflowed
- Fix overflowed node

$$(t-1) \leq \# \text{ of keys} \leq 2t-1$$

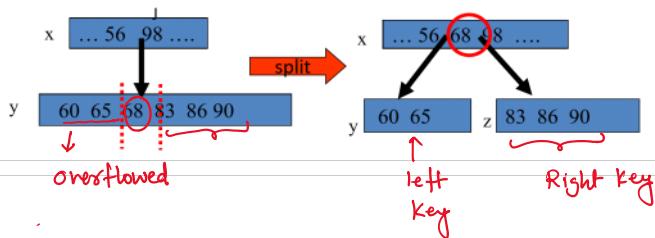
$$t \leq \# \text{ of children} \leq 2t$$

Fix an Overflowed

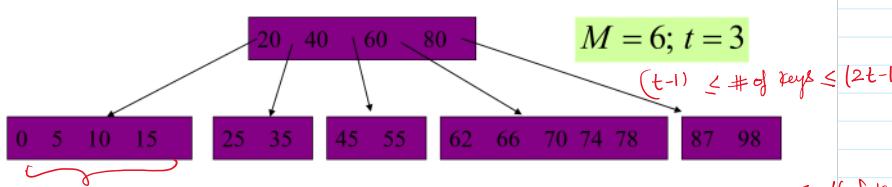
- Split the node into three parts, $M=2t$:
 - Left: the first t values, become a left child node
 - Middle: the middle value at position t , goes up to parent
 - Right: the last $t-1$ values, become a right child node
- Continue with the parent:
 - Until no overflow occurs in the parent
 - If the root overflows, split it too, and create a new root node

$(2t-1)$ key

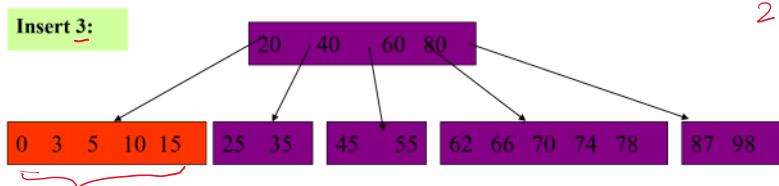
$2t$

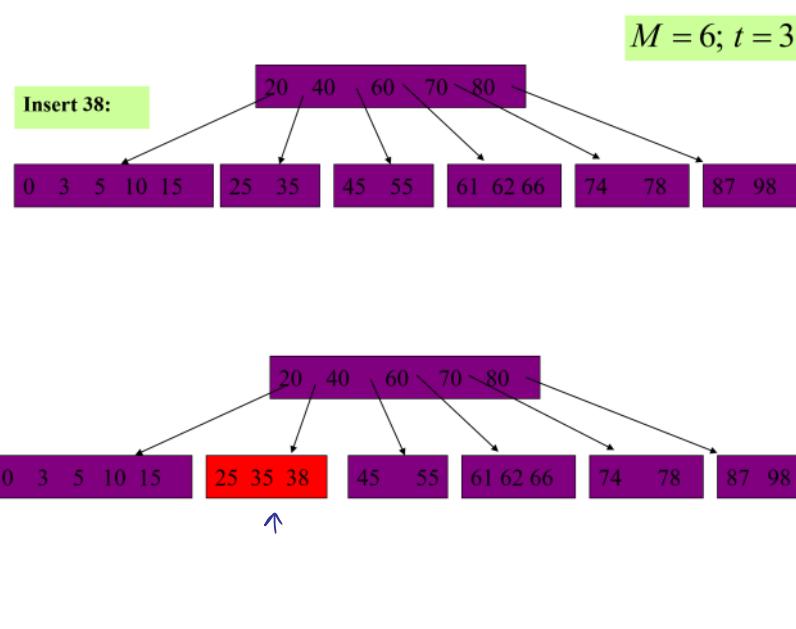
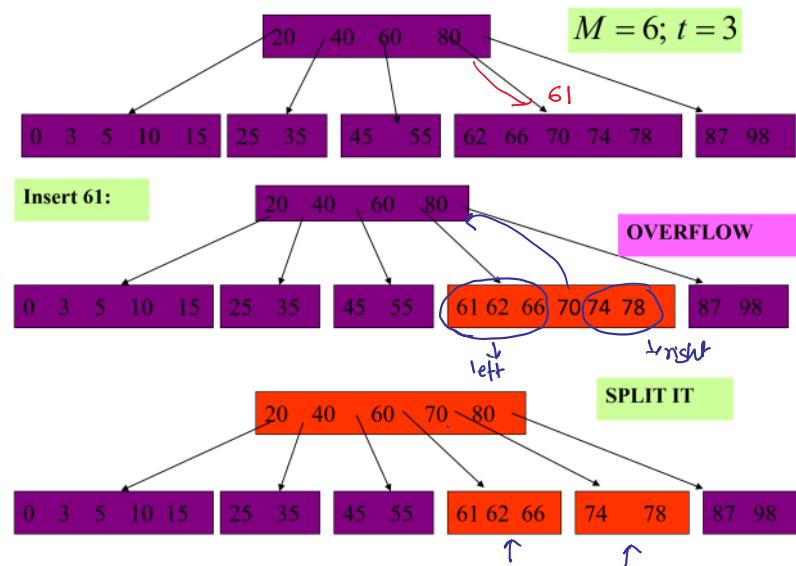


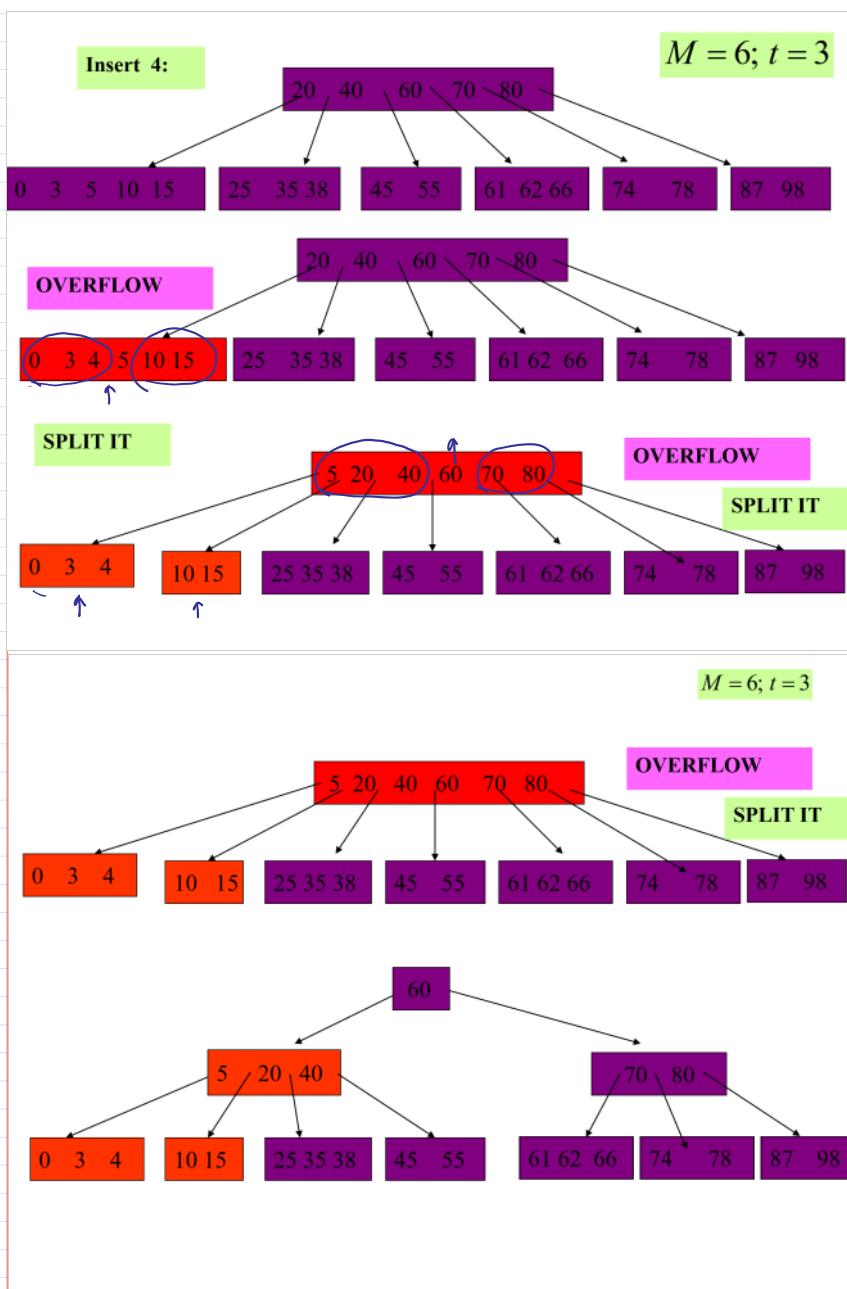
Insert example



$2 \leq \# \text{ of keys} \leq 5$







Complexity Insert

- Inserting a key into a B-tree of height h is done in a single pass down the tree and a single pass up the tree

Complexity: $O(h) = O(\log_t n)$

$$t-1 \leq \# \text{ keys} \leq 2t-1$$

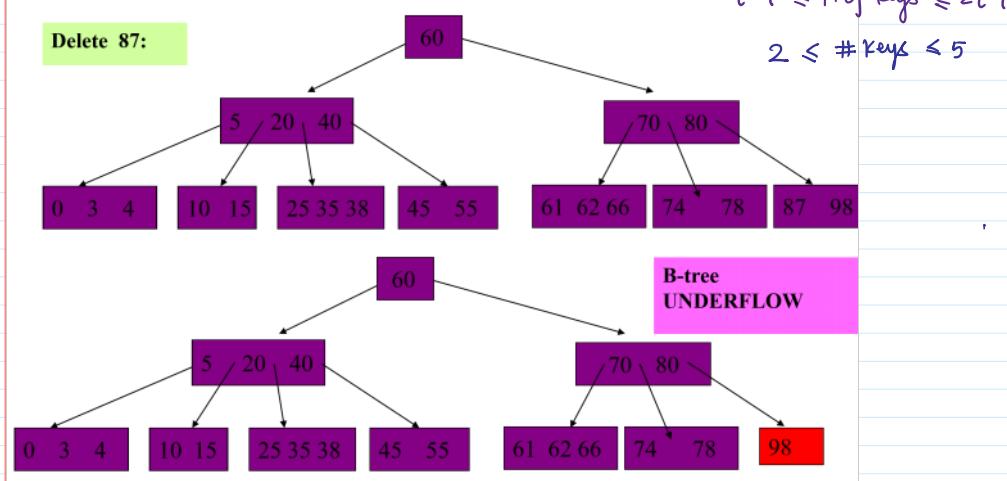
B-Tree: Delete X

- Delete as in M-way tree
- A problem:
 - might cause **underflow**: the number of keys remain in a node $< t-1$

Recall: The root should have at least 1 value in it, and all other nodes should have at least $t-1$ values in them

$M = 6; t = 3$

Underflow Example



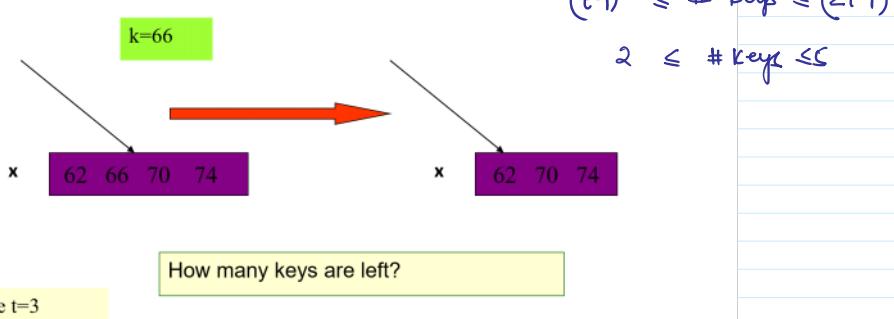
B-Tree: Delete X,k

- Delete as in M-way tree
- A problem:
 - might cause **underflow**: the number of keys remain in a node $< t-1$
- Solution:
 - make sure a node that is visited has at least t instead of $t-1$ keys.
 - If it doesn't have k
 - (1) either take from sibling via a rotate, or
 - (2) merge with the parent
 - If it does have k
 - See next slides

Recall: The root should have at least 1 value in it, and all other nodes should have at least $t-1$ (at most $2t-1$) values in them

B-Tree-Delete (x, k)

1st case: k is in x and x is a leaf \rightarrow delete k

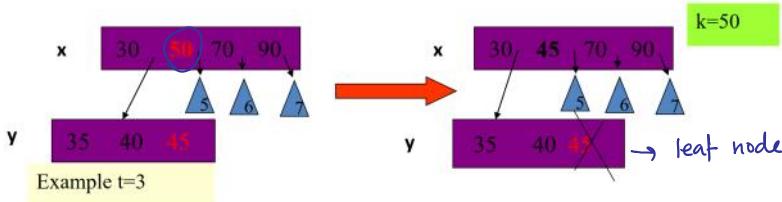


2nd case: k in the internal node x , y and z are the preceding and succeeding nodes of the key $k \in x$

a. If y has at least t keys:

- ▷ Replace k in x by $k' \in y$, where k' is the predecessor of k in y
- ▷ Delete k' recursively

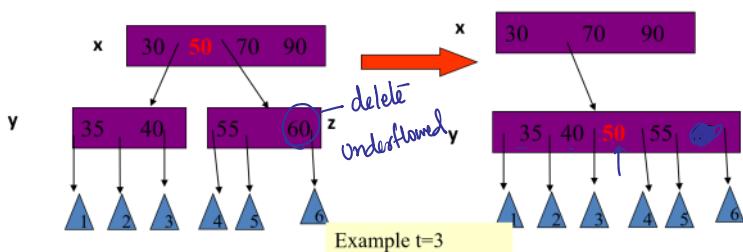
b. Similar check for successor case



2nd case cont.:

c. Both a and b are not satisfied: y and z have $t-1$ keys

- Merge the two children, y and z
- Recursively delete k from the merged cell



Questions

- When does the height of the tree shrink?
- Why do we need the number of keys to be at least t and not $t-1$ when we proceed down in the tree?

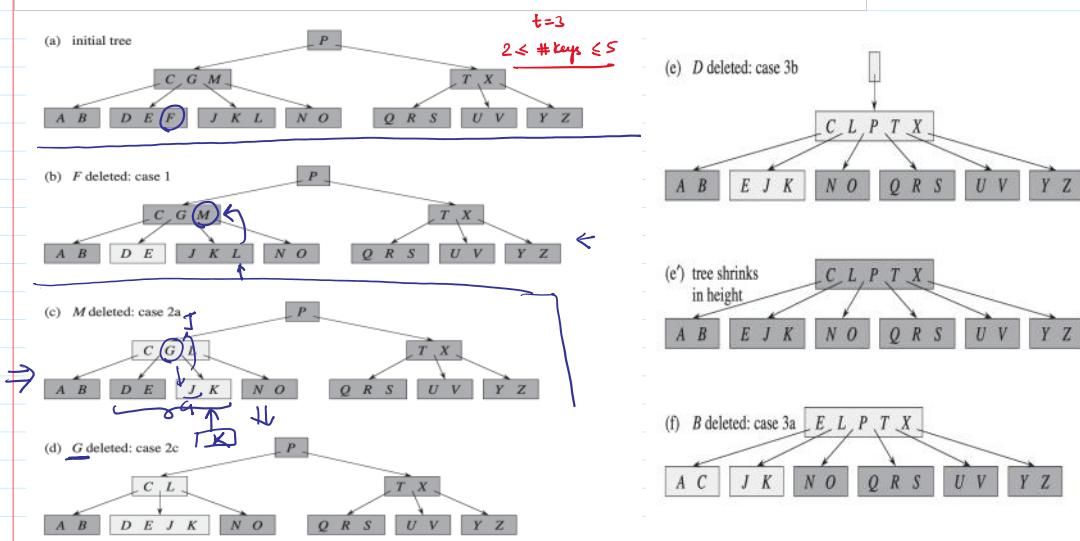


Figure 18.8 Deleting keys from a B-tree. The minimum degree for this B-tree is $t = 3$, so a node (other than the root) cannot have fewer than 2 keys. Nodes that are modified are lightly shaded.
 (a) The B-tree of Figure 18.7(e). (b) Deletion of F . This is case 1: simple deletion from a leaf.
 (c) Deletion of M . This is case 2a: the predecessor L of M moves up to take M 's position.
 (d) Deletion of G . This is case 2c: we push G down to make node $DEGJK$ and then delete G from this leaf (case 1).

Figure 18.8, continued (e) Deletion of D . This is case 3b: the recursion cannot descend to node CL because it has only 2 keys, so we push P down and merge it with CL and TX to form $CLPTX$; then we delete D from a leaf (case 1). (e') After (e), we delete the root and the tree shrinks in height by one. (f) Deletion of B . This is case 3a: C moves to fill B 's position and E moves to fill C 's position.

Delete Complexity

- Basically downward pass:
 - Most of the keys are in the leaves – one downward pass
 - When deleting a key in internal node – may have to go one step up to replace the key with its predecessor or successor

Complexity

$$O(h) = O(\log_t n)$$

Run Time Analysis of B-Tree Operations

- For a B-Tree of order $M=2t$
 - #keys in internal node: $M-1$
 - #children of internal node: between $M/2$ and M
 - Depth of B-Tree storing n items is $O(\log_{M/2} N)$
- Find run time is:
 - $O(\log M)$ to binary search which branch to take at each node, since M is constant it is $O(1)$.
 - Total time to find an item is $O(h * \log M) = O(\log n)$
- Insert & Delete
 - Similar to find but update a node may take : $O(M)=O(1)$

Note: if M is >32 it worth using binary search at each node

A typical B-Tree

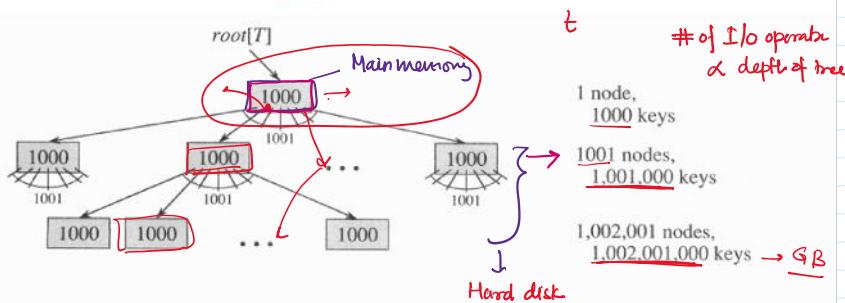
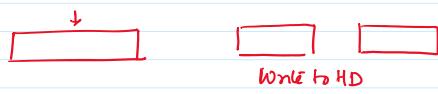


Figure 18.3 A B-tree of height 2 containing over one billion keys. Each internal node and leaf contains 1000 keys. There are 1001 nodes at depth 1 and over one million leaves at depth 2. Shown inside each node x is $n[x]$, the number of keys in x .

Copyright © The McGraw-Hill Companies, Inc. Permission required for reproduction or display.



Why B-Tree?

- B-trees is an implementation of dynamic sets that is optimized for disks
 - The memory has an hierarchy and there is a tradeoff between size of units/blocks and access time
 - The goal is to optimize the number of times needed to access an “expensive access time memory”
 - The size of a node is determined by characteristics of the disk – block size – page size
 - The number of access is proportional to the tree depth

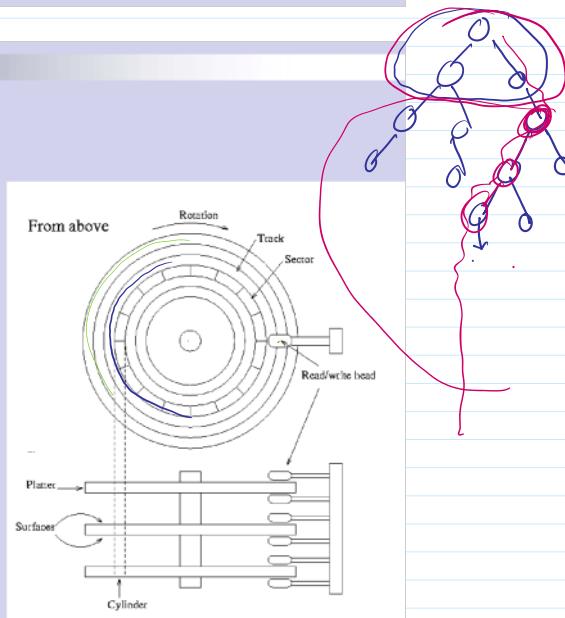
Disk Based Data Structures

- So far search trees were limited to main memory structures
 - Assumption: the dataset organized in a search tree fits in main memory (including the tree overhead)
- Counter-example: transaction data of a bank > 1 GB per day
 - use secondary storage media (punch cards, hard disks, magnetic tapes, etc.)
- Consequence: make a search tree structure secondary-storage-enabled

Key value

Hard Disks

- Large amounts of storage, but slow access!
- Identifying a page takes a long time (seek time plus rotational delay – 5-10ms), reading it is fast
 - It pays off to read or write data in **pages** (or blocks) of 2-16 Kb in size.



Algorithm analysis

- The running time of disk-based algorithms is measured in terms of
 - computing time (CPU)
 - number of disk accesses
 - sequential reads
 - random reads
- Regular main-memory algorithms that work one data element at a time can not be “ported” to secondary storage in a straight-forward way